

Διάλεξη #1 - Security Fundamentals

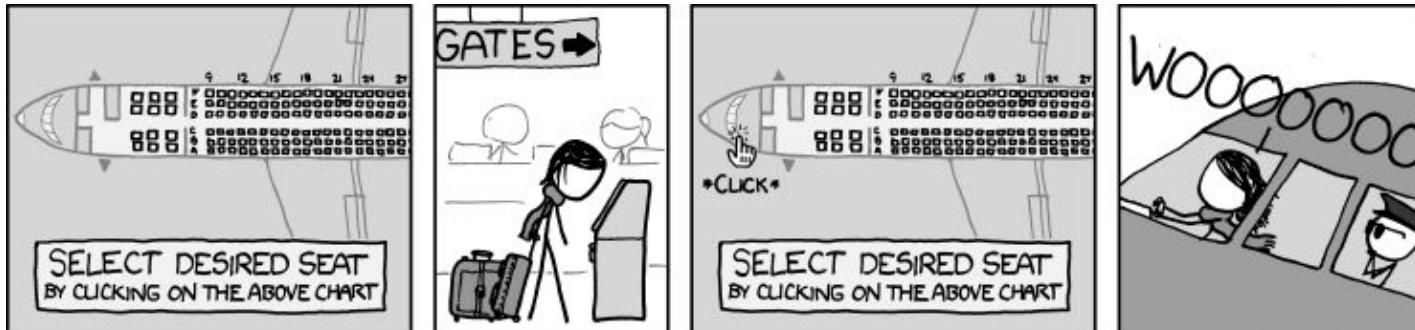
Εθνικό και Καποδιστριακό Πανεπιστήμιο Αθηνών

Εισαγωγή στην Ασφάλεια

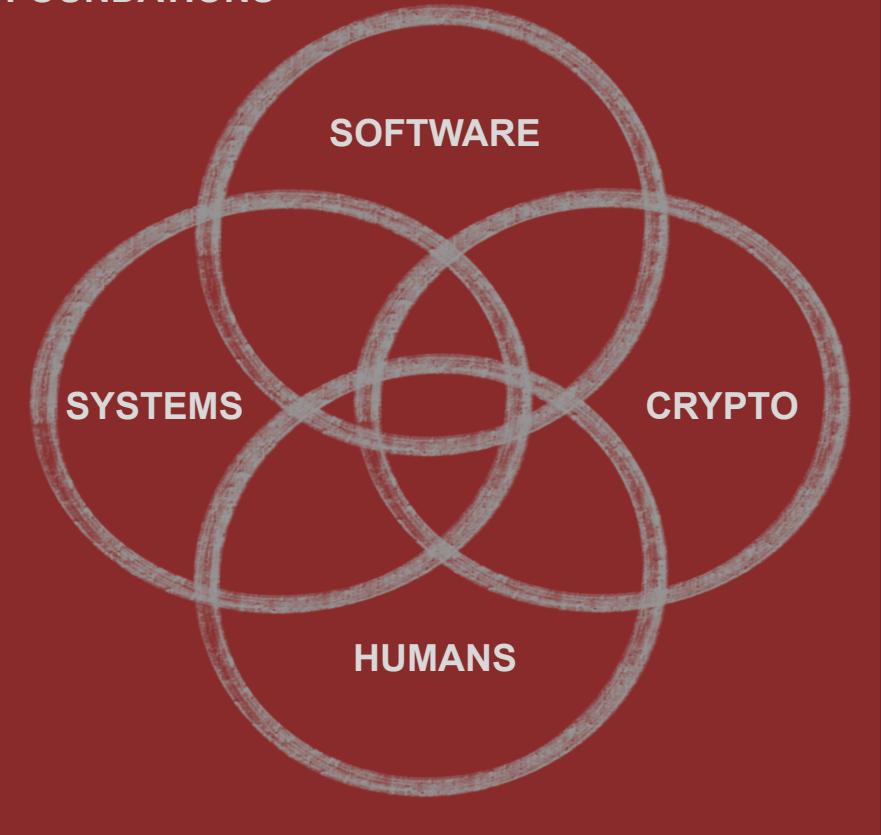
Θανάσης Αυγερινός

(Security Mindset)

<https://xkcd.com/726/>



FOUNDATIONS



Huge thank you to [David Brumley](#) from Carnegie Mellon University for the guidance and content input while developing this class

Ανακοινώσεις / Διευκρινίσεις

- Άρα δεν επιτρέπεται να κάνουμε `scanf("%s", buffer);`

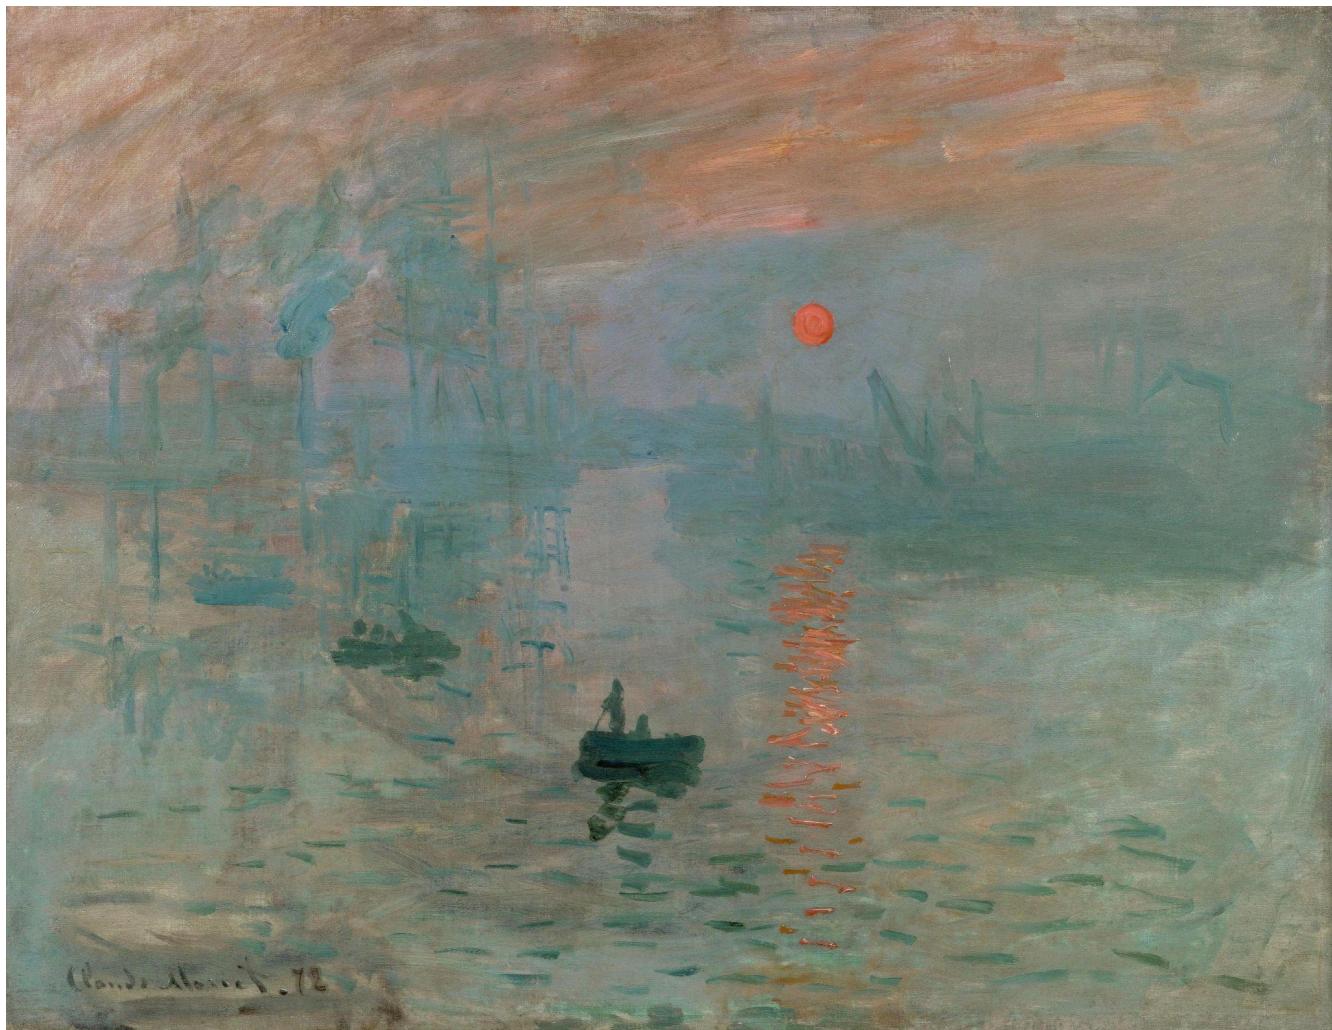
Την Προηγούμενη Φορά

1. Διαδικαστικά
2. Σκοπός του μαθήματος
3. Ασφάλεια και Συστήματα
4. Σχέδιο για το μάθημα φέτος
5. Το πρώτο μας exploit



Σήμερα

- Security Fundamentals
 - Adversaries
 - Threat Models
 - Security Properties
 - Trusted Computing Base (TCB)
 - Security Principles



Two Concepts
(Only a few I need
you to memorize)

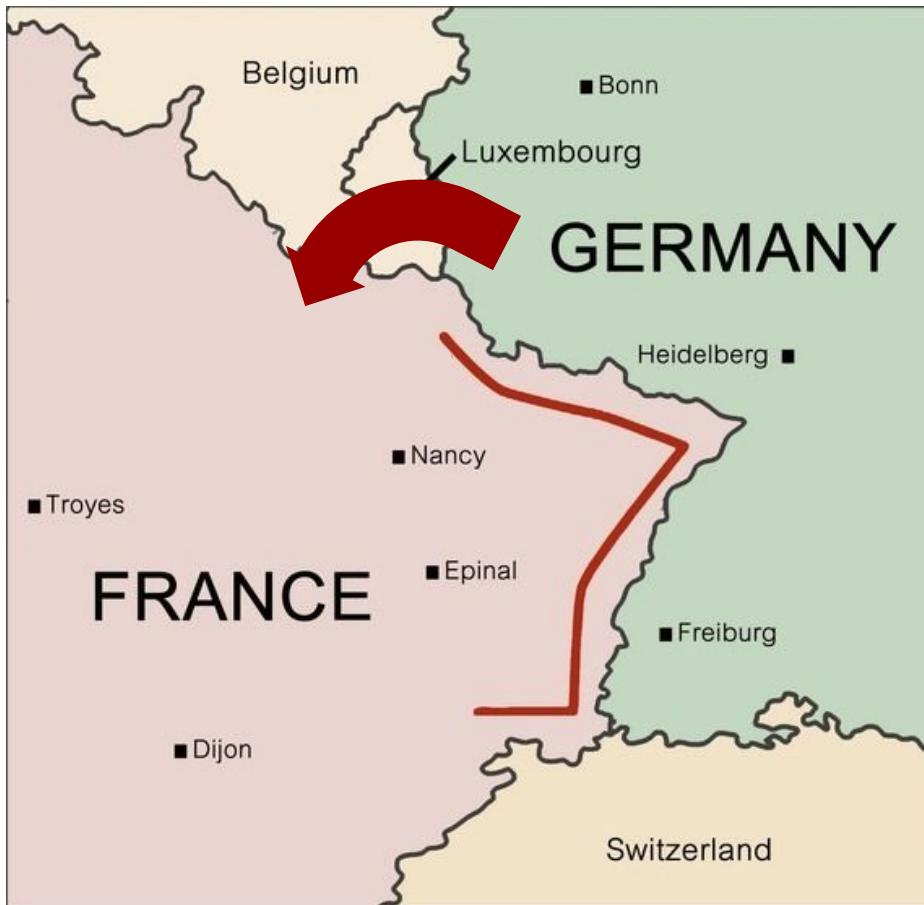
Defining The Adversary (1/2)

- Adversary = < Goals, Capabilities >
- Goal: What constitutes success?
 - May involve subgoals
 - Example goal: Gain control of X's data
 - Sub-goal: Reconnaissance: search online for info about X
 - Sub-goal: Access: Guess X's ssh password on Linux lab
 - » Sub-goal: Lateral movement: Use ssh account to move to other services / linked accounts
- Capabilities: What resources can the adversary use?
 - 1 computer or millions?
 - Physical or remote access?
 - Access to source code?

Why don't we include
adversary's strategy?

Defining The Adversary (1/2)

- Adversary = < Goals, Capabilities >

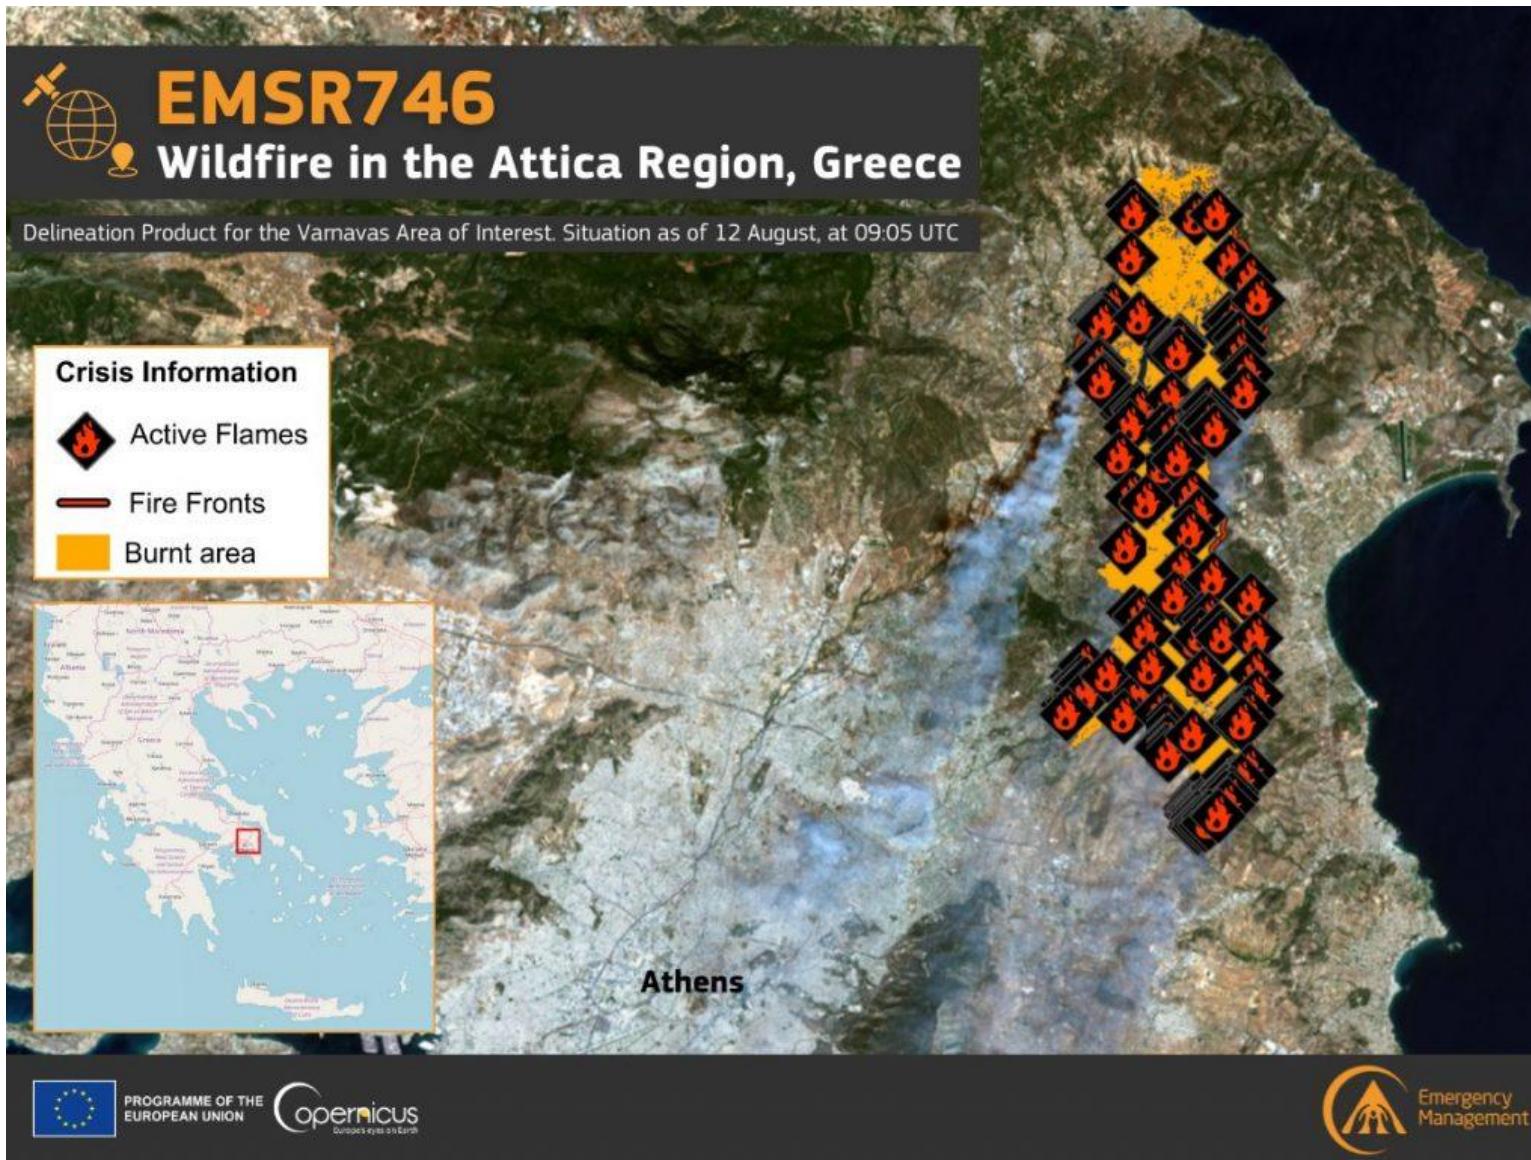


Why don't we include
adversary's strategy?

Security Mechanism Classification for a property (2/2)

1. **Prevention.** Prevent issues from happening. Any precautionary measures.
2. **Detection.** Assuming an incident took place, detect them as early and as accurately as possible.
3. **Resilience.** Assuming one or multiple incidents took place, ensure the overall system security degrades gracefully and does not collapse.
4. **Deterrence.** Measures to ensure penalties for actors responsible for security incidents. Policy-based.

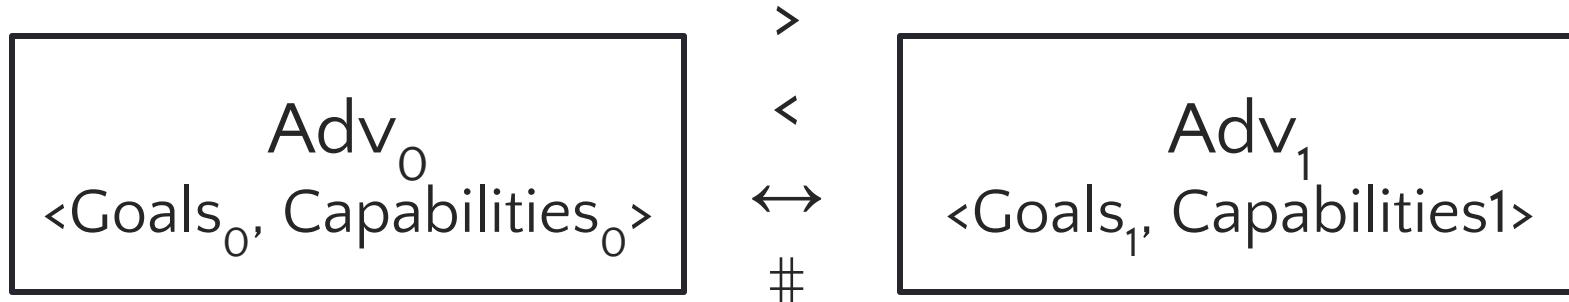
Fire Drill: How to Defend?



Adversaries

Basic Adversary Metrics

Partial Orders



Implication ($A_0 \rightarrow A_1$):

A_0 's goals and capabilities are a superset of A_1 's

Separation ($A_1 \not\rightarrow A_0$):

A_1 's goals and/or capabilities are not a superset of A_0 's

Strict Dominance $A_0 > A_1 = A_0 \rightarrow A_1 \wedge A_1 \not\rightarrow A_0$

Equivalence: $A_0 \leftrightarrow A_1 = A_0 \rightarrow A_1 \wedge A_1 \rightarrow A_0$

Incomparable: $A_0 \# A_1 = A_0 \not\rightarrow A_1 \wedge A_1 \not\rightarrow A_0$

Example Adversary Comparison

Adversary A0

- Goal: Modify the my-uni website
- Capabilities
 - View the website
 - Send data to the website
 - Modify local browser state
 - Access a normal user account on my-uni
 - Invoke system calls on my-uni.uoa.gr

Adversary A1

- Goal: Modify the my-uni website
- Capabilities
 - View the website
 - Send data to the website
 - Modify local browser state
 - ~~– Access a normal user account on
my studies~~
 - ~~– Invoke system calls on my-uni.uoa.gr~~

$$A0 \rightarrow A1 \wedge A1 \not\rightarrow A0$$

hence $A0 > A1$, i.e., A0 strictly dominates A1

Threat Models

Threat Modeling

Threat modeling is a process by which potential threats, such as [structural vulnerabilities](#) or the absence of appropriate safeguards, can be identified and enumerated, and countermeasures prioritized.

https://en.wikipedia.org/wiki/Threat_model

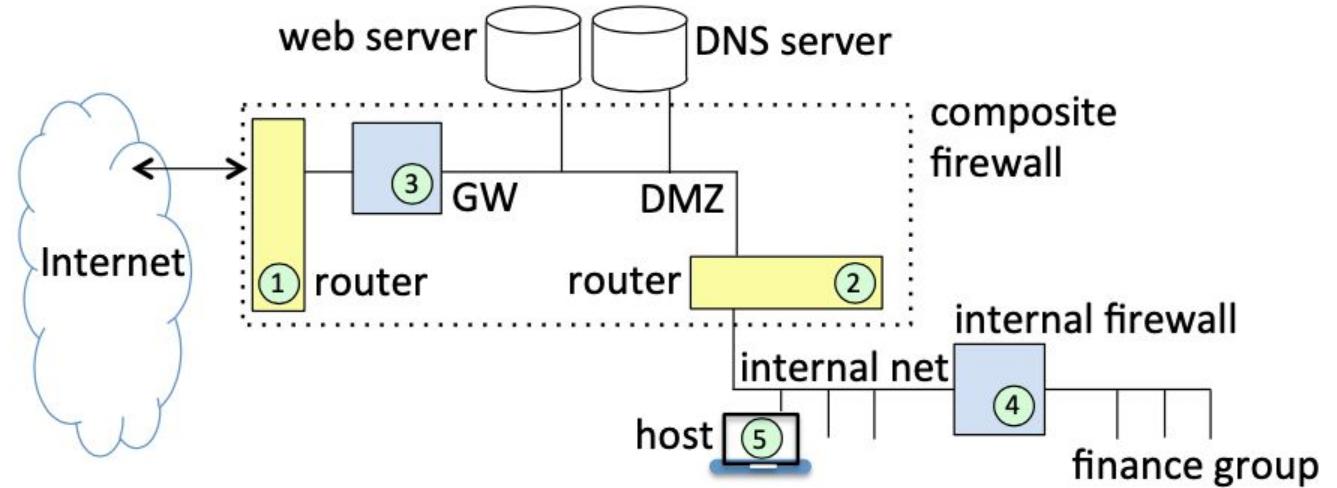
<https://www.threatmodelingmanifesto.org/>

Threat Model Includes

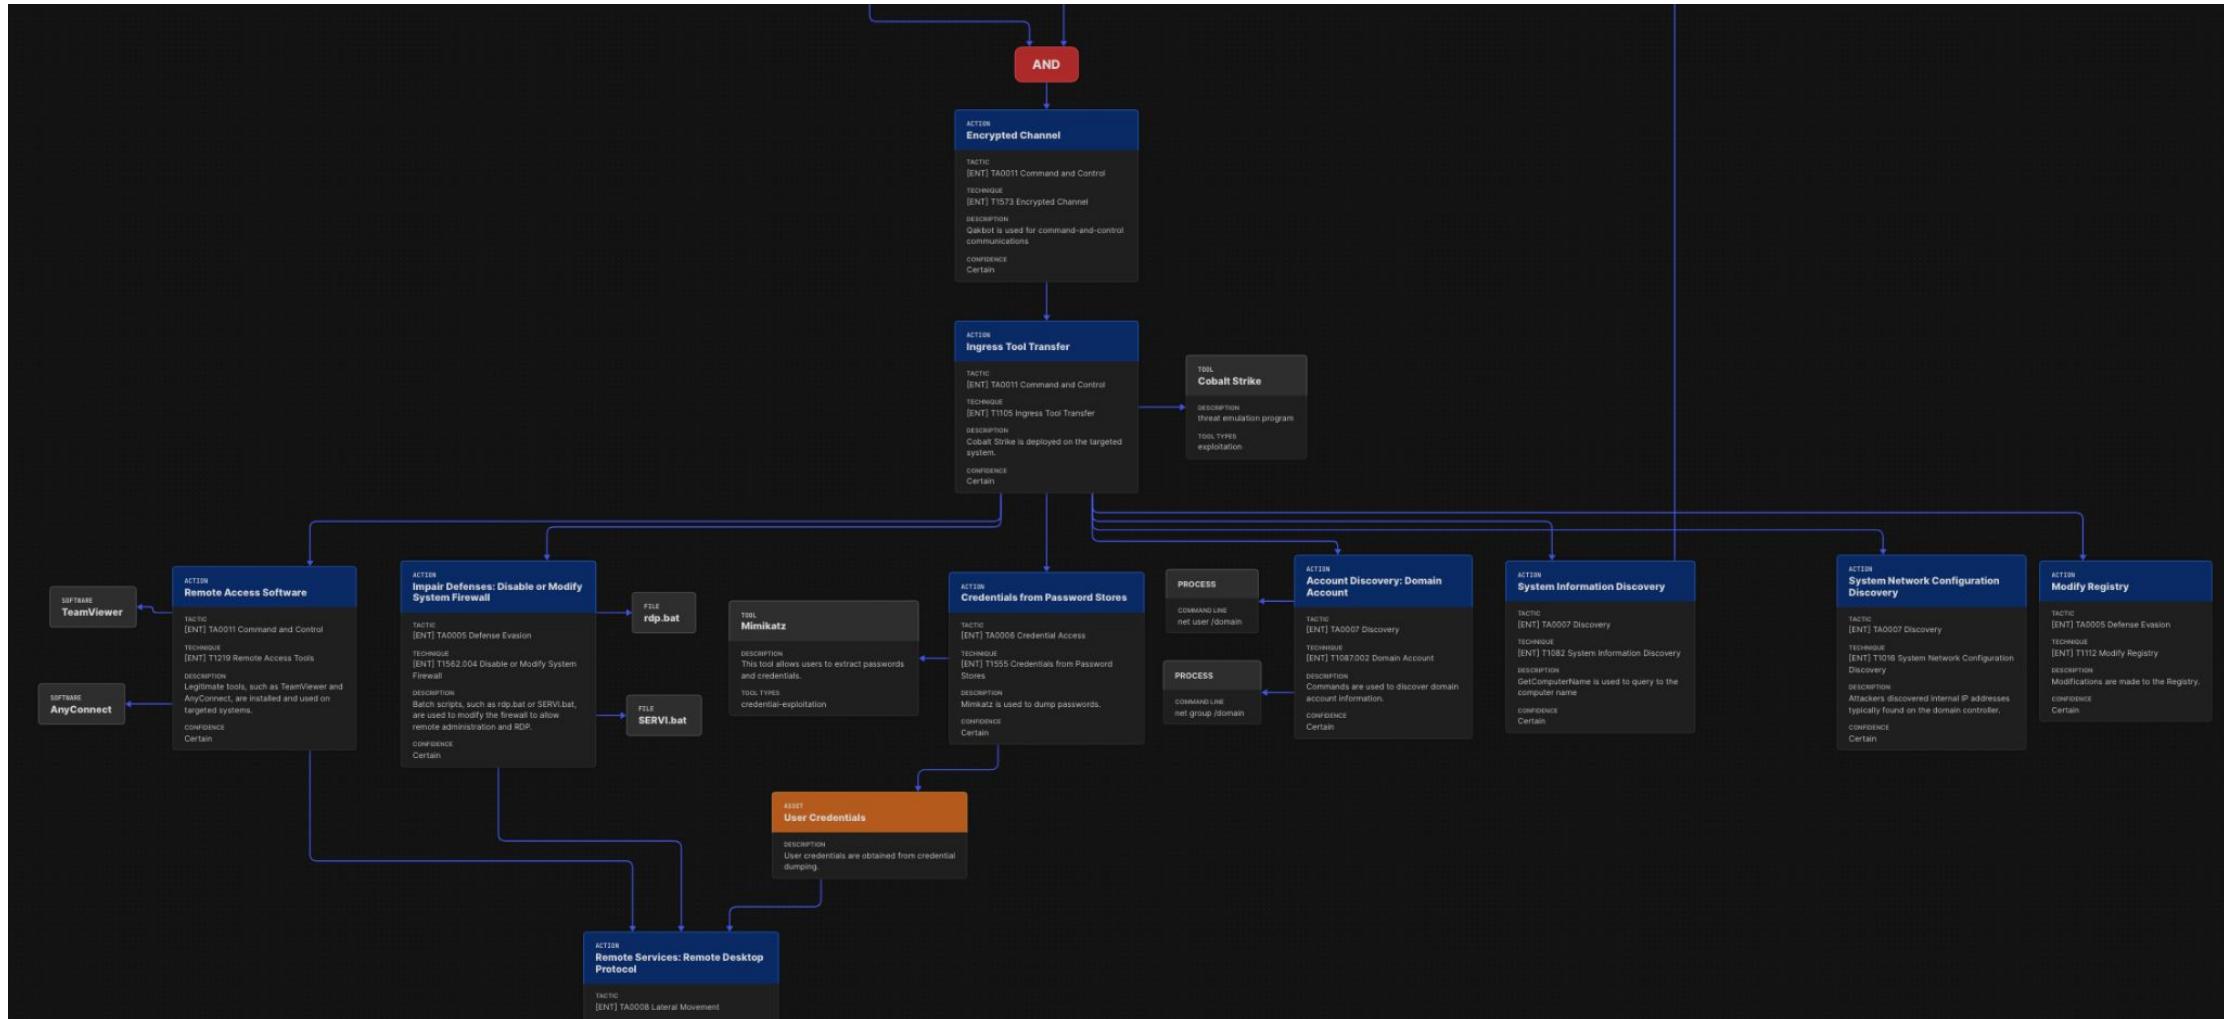
- Assets
 - What are you protecting?
 - Which matter most?
- System goals (functionality, security)
- Adversary definition – key characteristic
 - Risk assessment – we are in the insurance business!
 - Risk justifies the cost

Systematic Threat Modeling

- Diagram-based
- Attack trees
- Checklists
- STRIDE
- [MITRE ATT&CK](#)
- [Tactics, Techniques and Procedures \(TTP\)](#)



Attack Formalizations & Visualizations



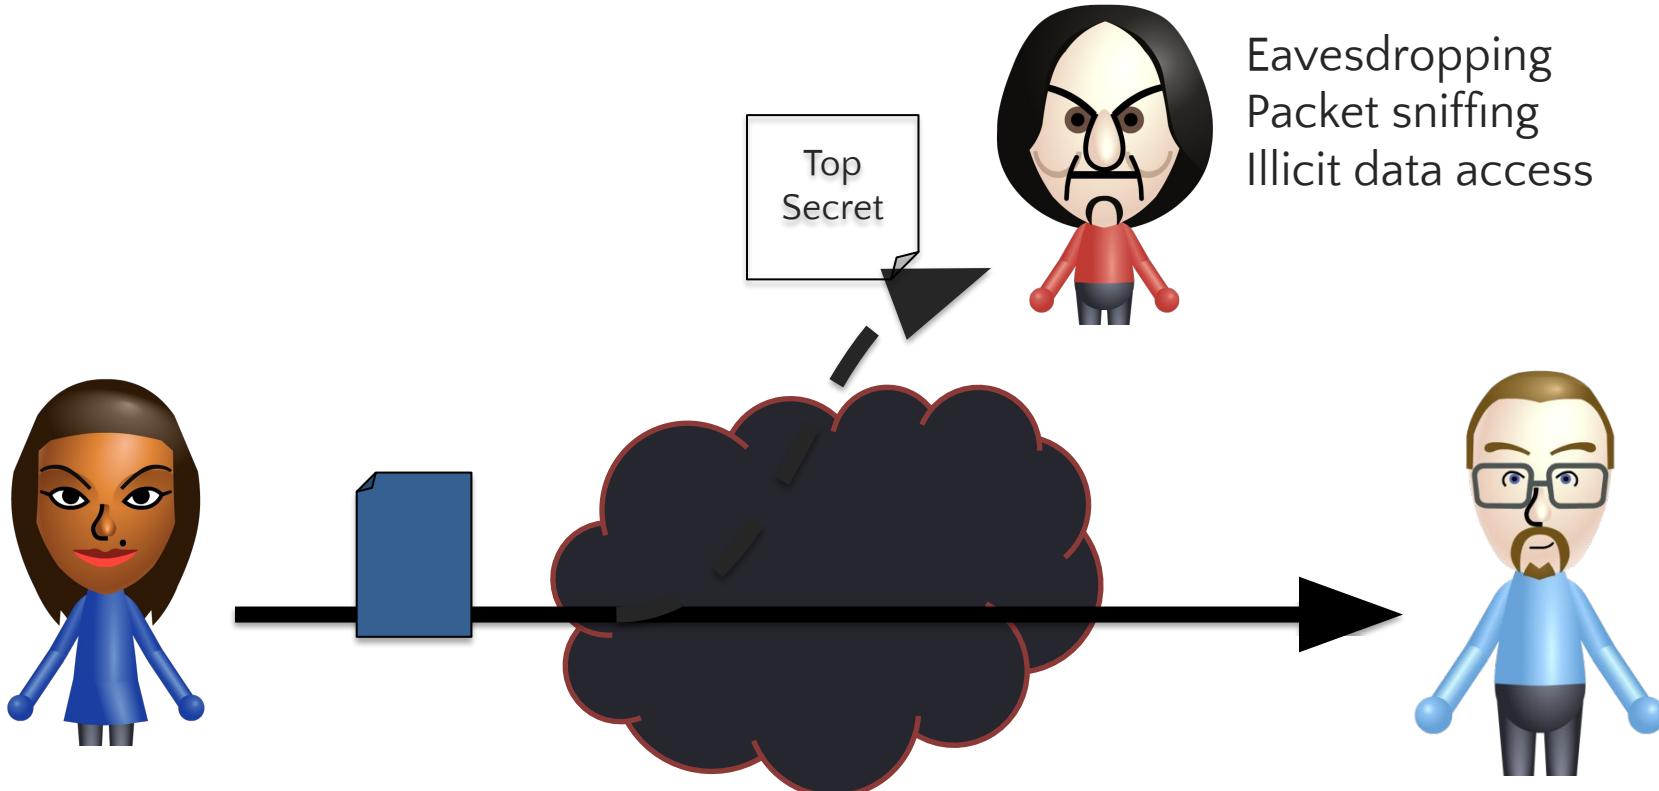
https://center-for-threat-informed-defense.github.io/attack-flow/example_flows/

<https://oasis-open.github.io/cti-documentation/stix/intro.html>

4 Key Security Properties

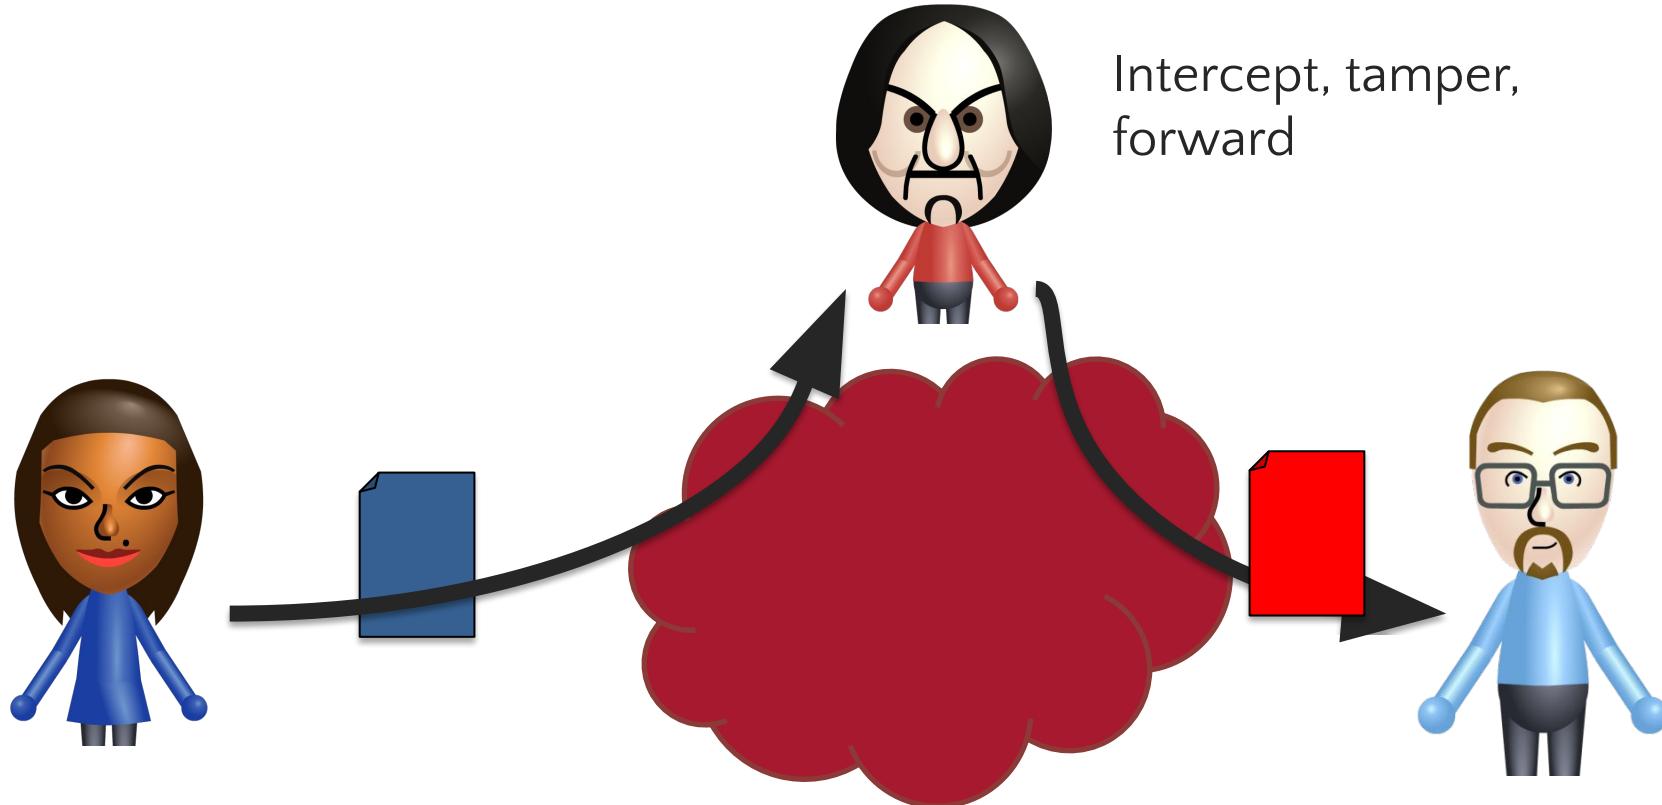
What Properties Do You Care About?

(1) Confidentiality == Secrecy == Concealing information



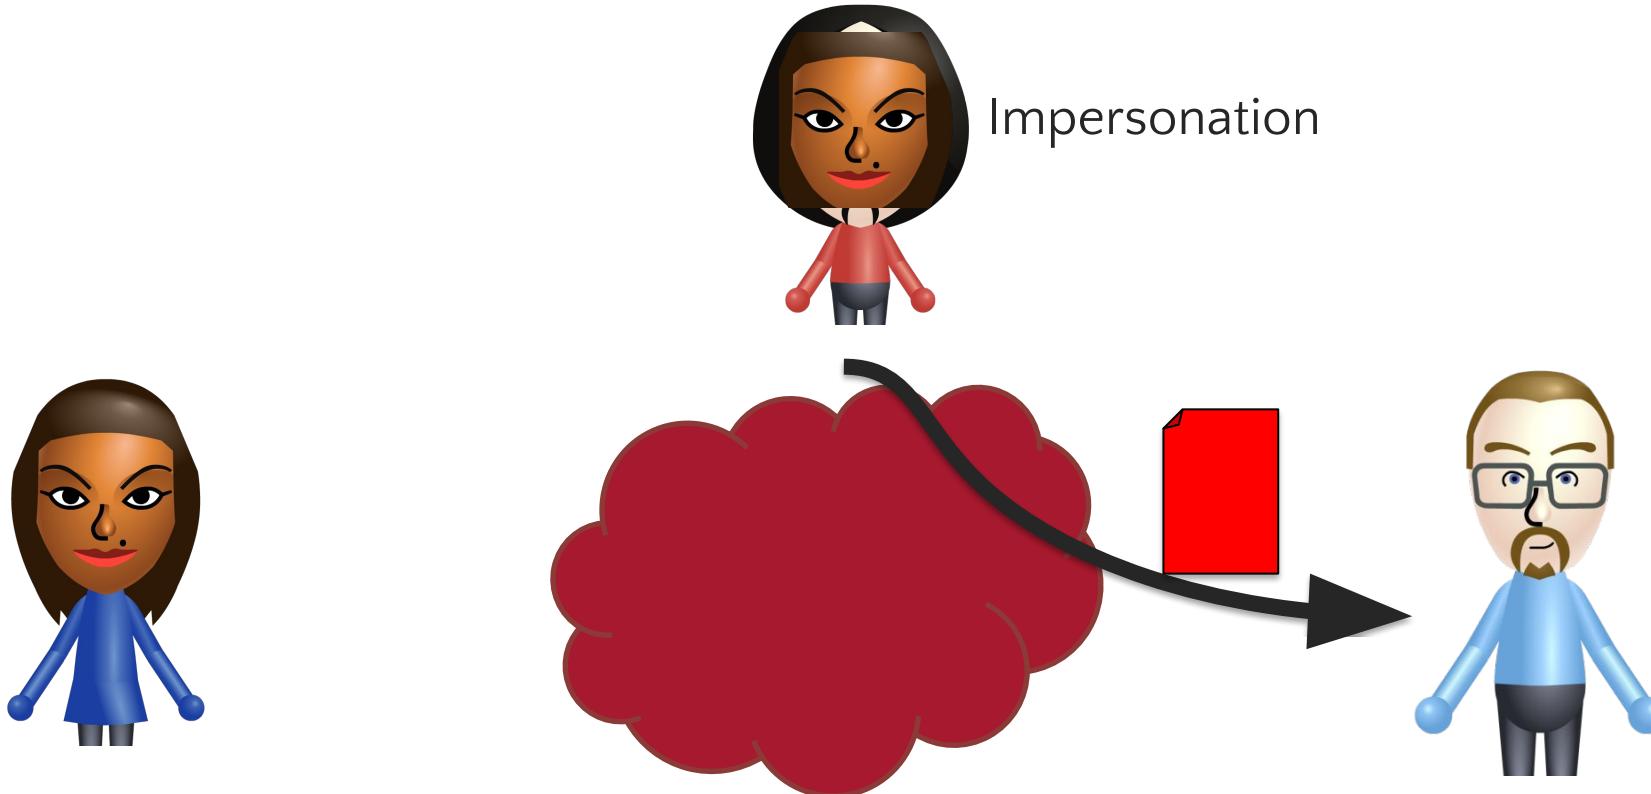
What Properties Do You Care About?

(2) **Integrity** == Prevention of unauthorized changes



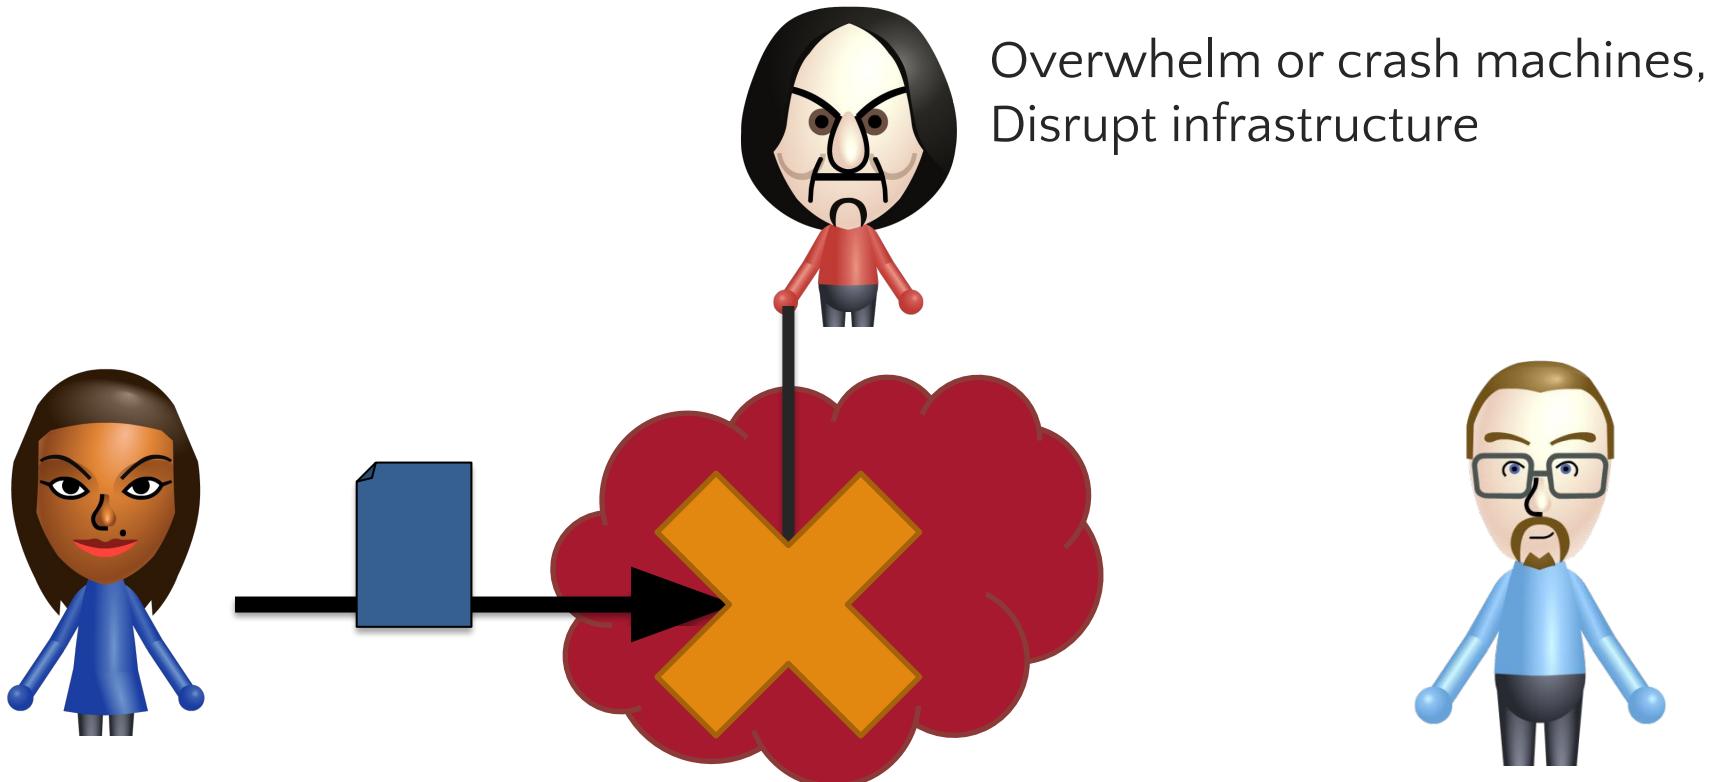
What Properties Do You Care About?

(3) **Authenticity** == Data and actions attributed to correct person



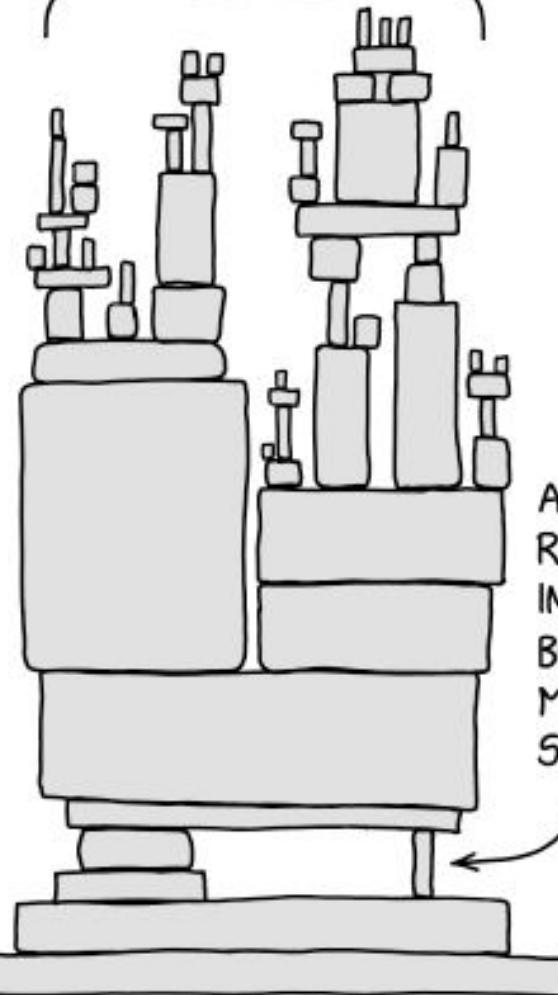
What Properties Do You Care About?

(4) Availability == Ability to use resources when needed



Trusted Computing Base (TCB)

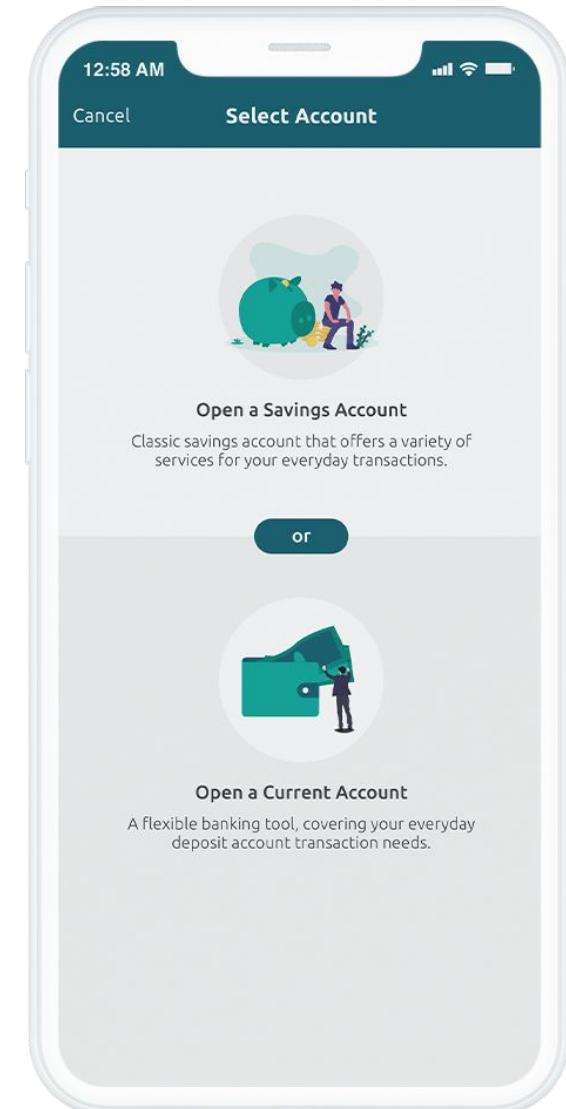
ALL MODERN DIGITAL
INFRASTRUCTURE



<https://xkcd.com/2347/>

Why Should I Trust This?

What does this application
rely on for its security?



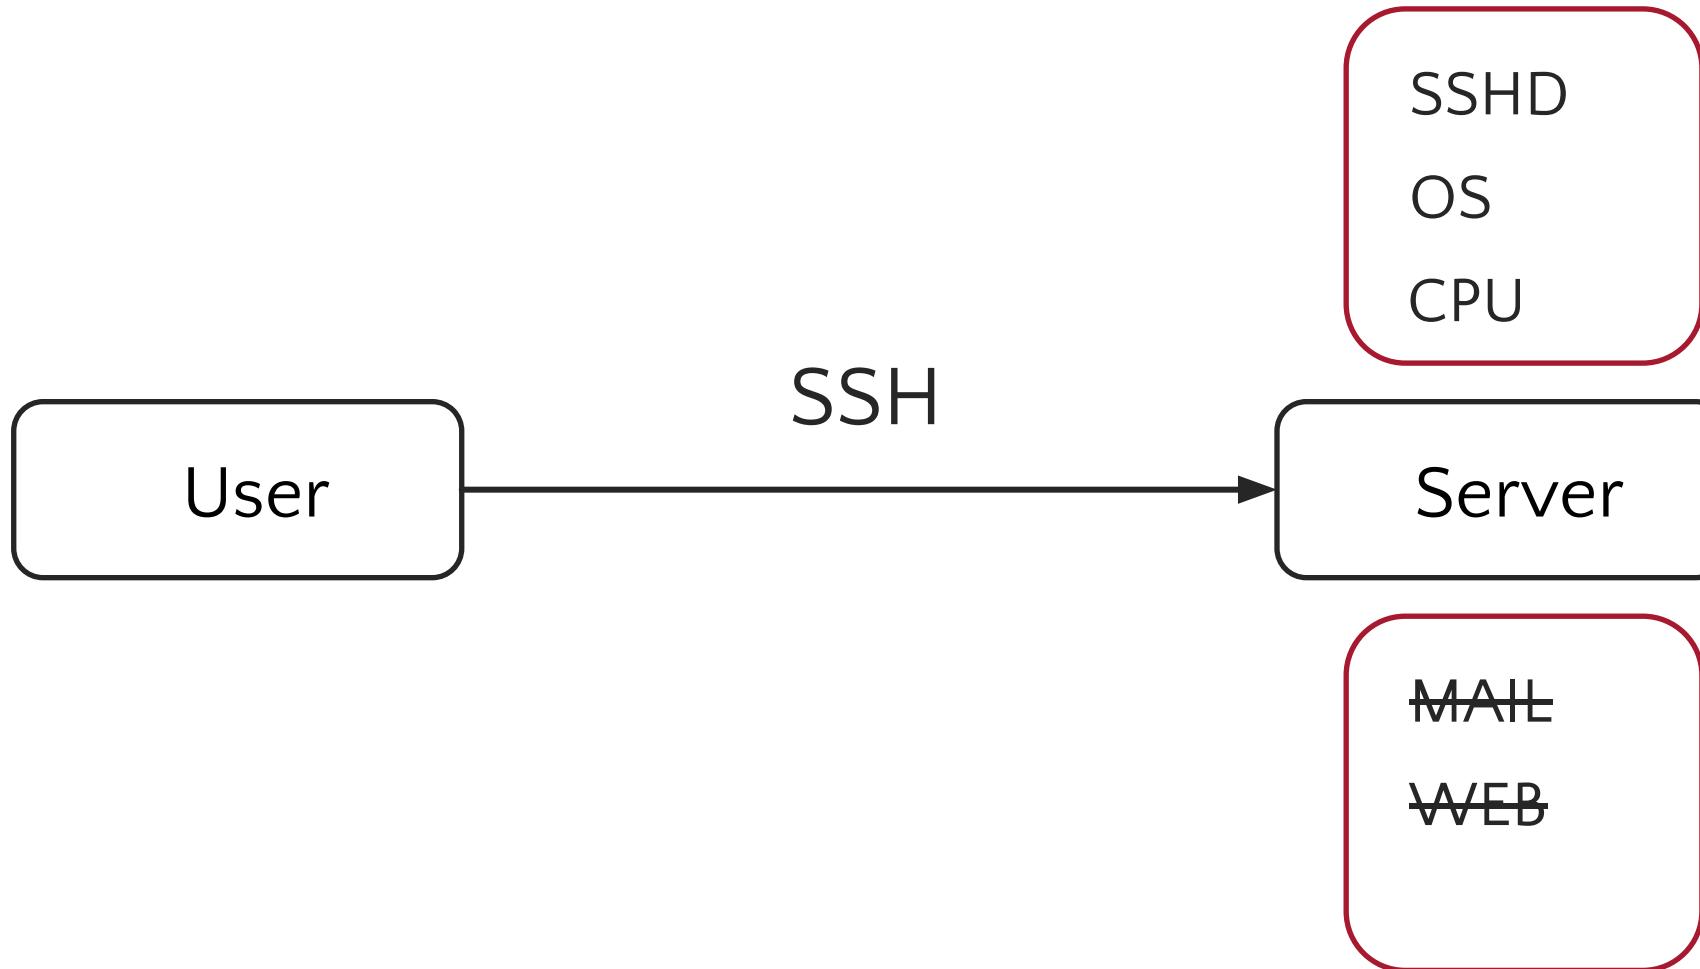
Trusted Computing Base (TCB)

- Component X's TCB is all other components that must operate securely for X to be secure
- Corollary 1: If TCB is secure, X has a chance of being secure
- Corollary 2: If TCB misbehaves, no guarantees about X's security!
- Trusted != Trustworthy

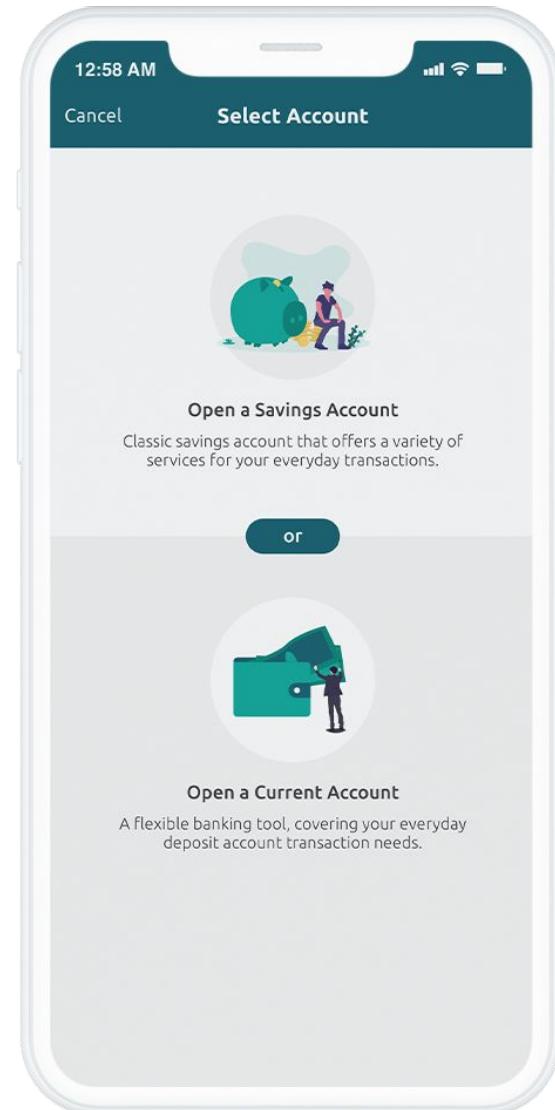
Example of TCB



Example of TCB



What is the TCB here?



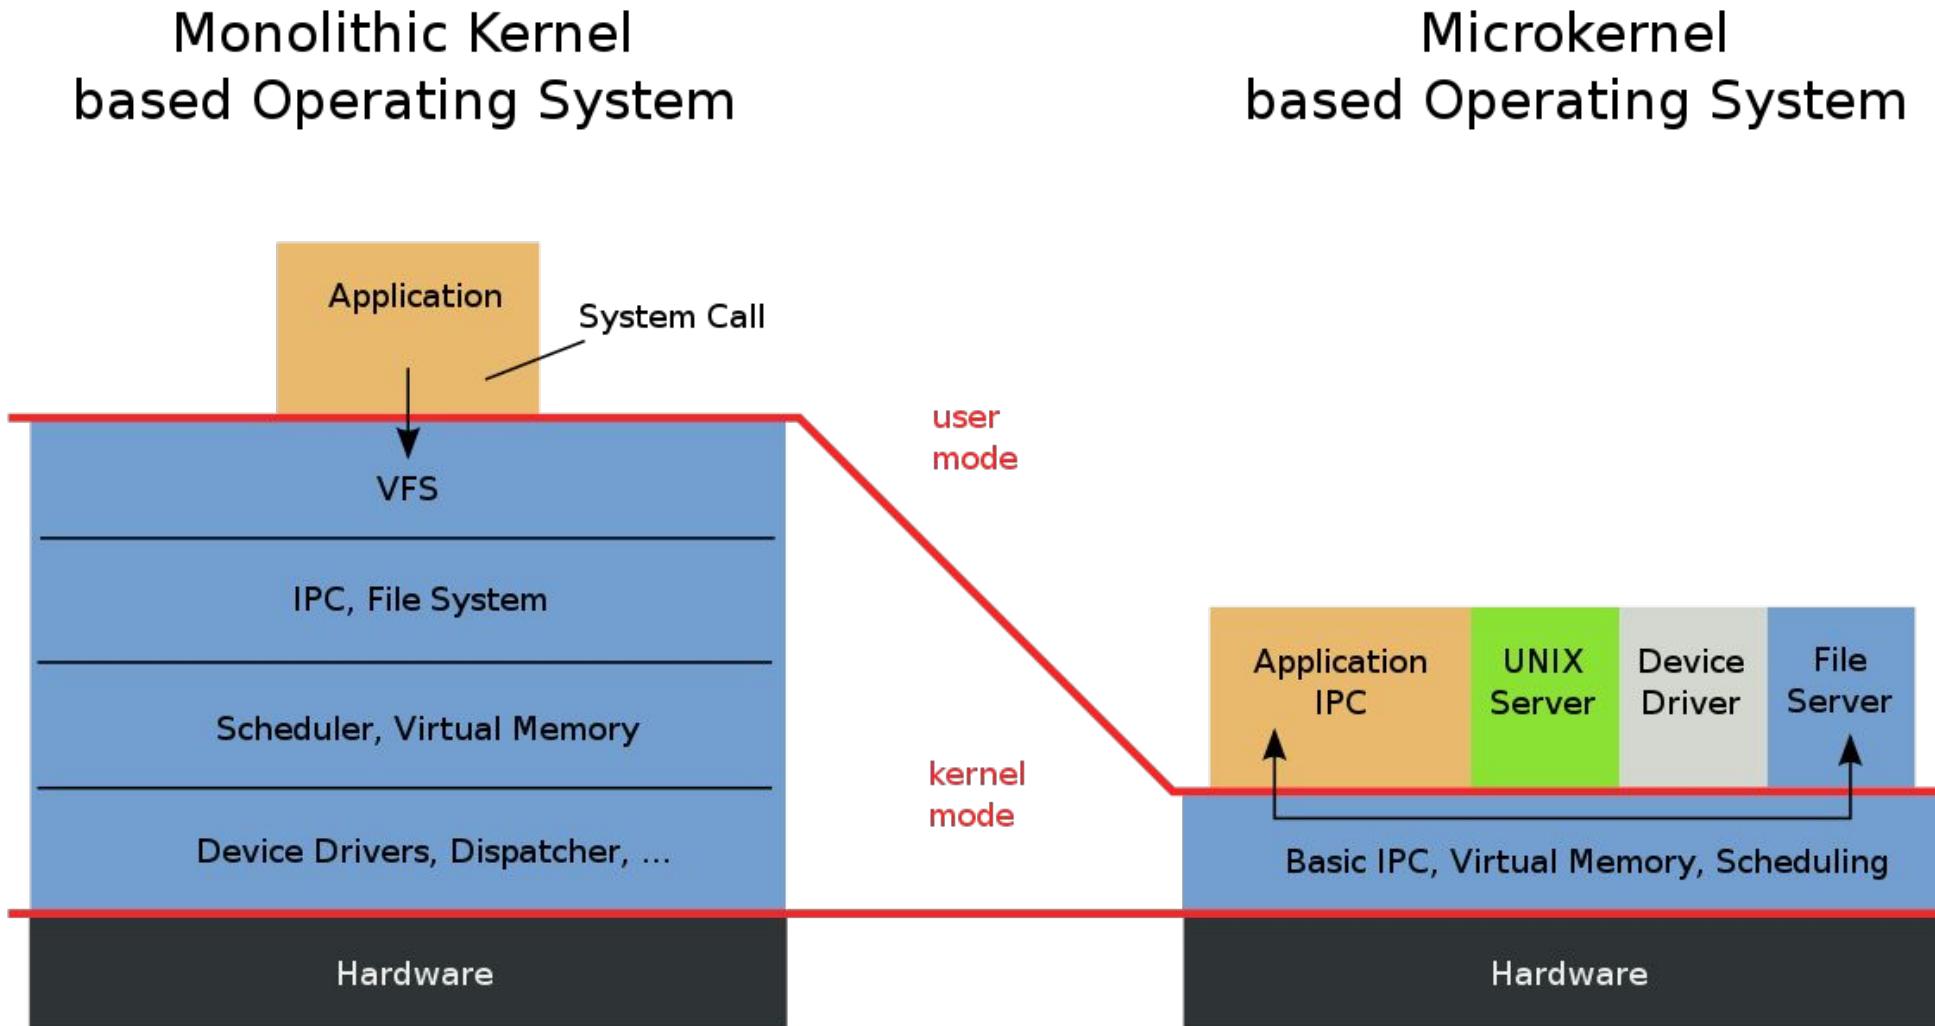
Ideal TCB Design

- Verifiable
 - Implies you want TCB to be as small as possible
 - (even when not verifiable, smaller = less buggy)
- Tamper proof
 - E.g., must prevent messing with the SSHD or OS executables

Why Do We Care About a TCB?

- Securing every piece of a system is hard!
- Identifying the TCB allows us to separate a system into a part that **must** be trusted and a part that **doesn't have to be**
- Can focus security efforts on the trusted piece
 - Reason about security more rigorously
- Caveat: Determining TCB is easier said than done

Example: Operating system kernel



Participation Question

Which of the following is *NOT* in the TCB of a *web browser* on your laptop?

- A. The laptop's OS
- B. JavaScript the browser downloads when you visit a website
- C. The laptop's hardware
- D. The browser's cryptographic library

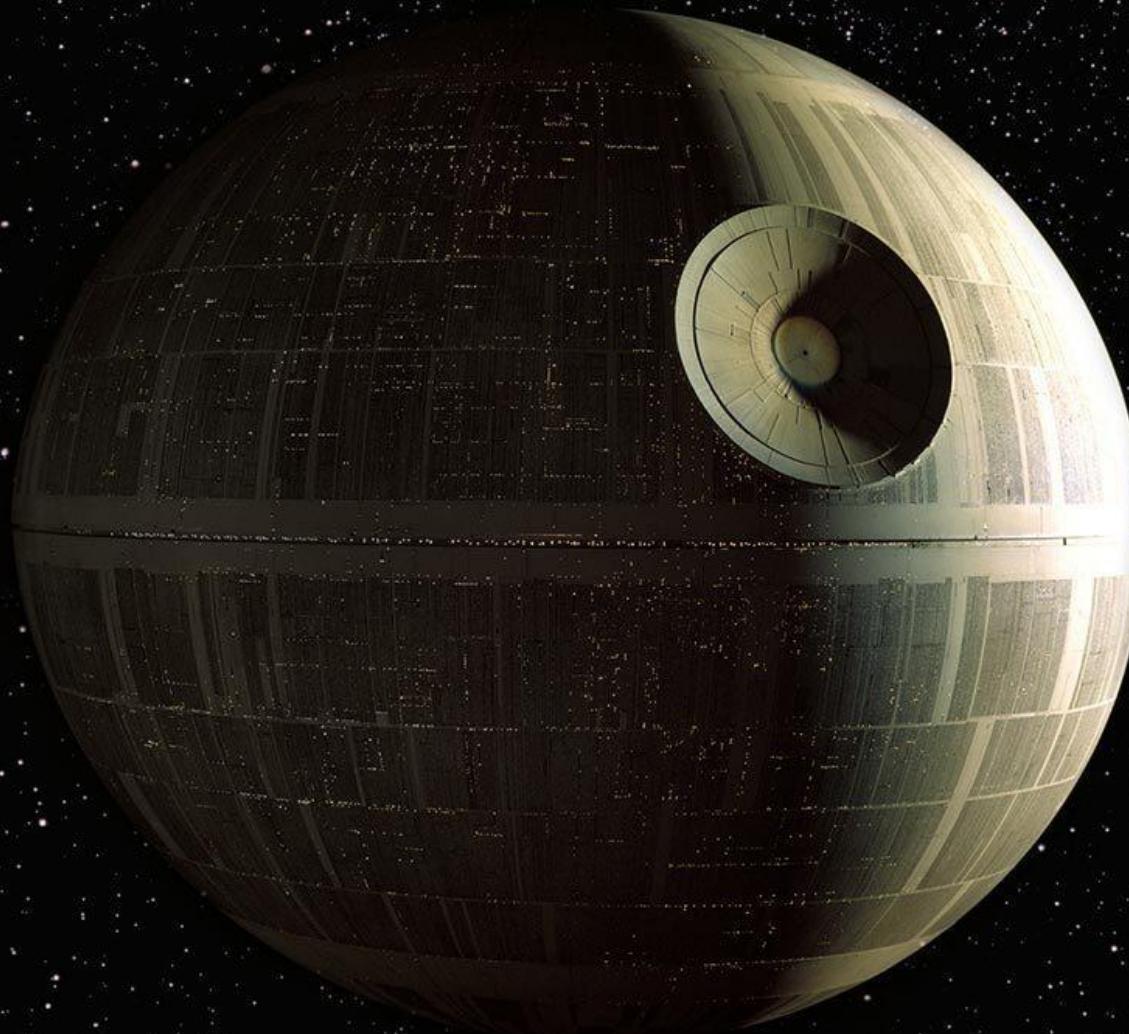
Designing Secure Systems

The Key Principles

- Economy of mechanism a.k.a. KISS
- Fail-safe defaults
- Don't rely on security by obscurity
- Complete mediation
- Least privilege
- Separation of duty
- Defense in depth
- Factor in users/acceptance/psychology
- Work factor/economics

} Later

See the reading for more useful principles



Keep It Simple, Stupid (KISS)

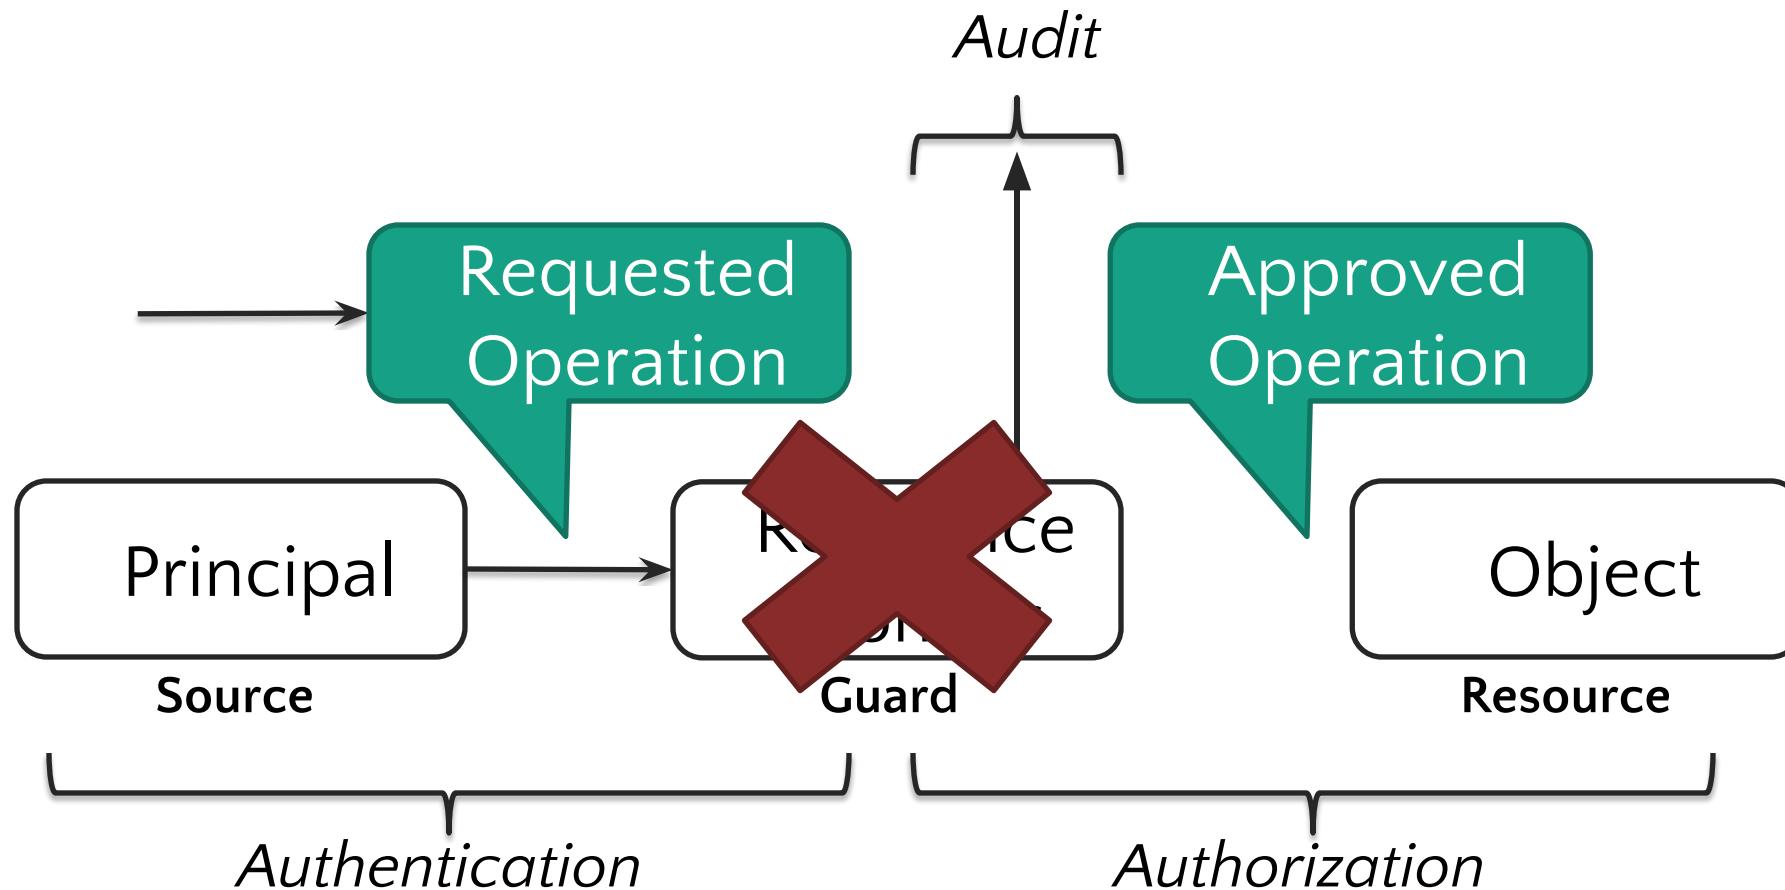
- Rule of thumb:
1-5 defects per 1K lines of code
- Windows 10 = 50M LOC; Linux 6.7 = 27M LOC
 - In both cases, essentially all in the TCB
- Smaller, simpler TCB is easier to reason about
 - e.g.: seL4 (a formally verified microkernel) = 89k LOC

The Key Principles

- Economy of mechanism a.k.a. KISS
- **Fail-safe defaults**
- Don't rely on security by obscurity
- Complete mediation
- Least privilege
- Separation of duty
- Defense in depth
- Factor in users/acceptance/psychology
- Work factor/economics

} Later

Fail-safe Defaults (Fail Closed)



The Key Principles

- Economy of mechanism a.k.a. KISS
- Fail-safe defaults
- **Don't rely on security by obscurity**
- Complete mediation
- Least privilege
- Separation of duty
- Defense in depth
- Factor in users/acceptance/psychology
- Work factor/economics

} Later

No Security by Obscurity



<https://xkcd.com/257/>

The Key Principles

- Economy of mechanism a.k.a. KISS
- Fail-safe defaults
- Don't rely on security by obscurity
- **Complete mediation**
- Least privilege
- Separation of duty
- Defense in depth
- Factor in users/acceptance/psychology
- Work factor/economics

} Later

Complete Mediation

- Every access to every object is checked by the reference monitor
- Easier said than done!
- TOCTTOU problems



Mediation: TOCTTTOU Vulnerabilities

Time-Of-Check-To-Time-Of-Use

```
int openfile(char *path){  
    struct stat s;  
    if (stat(path,&s) < 0))  
        return -1;  
  
    if (!S_ISREG(s.st_mode)){  
        error("only regular files allowed");  
        return -1;  
    }  
    return open(path,0_RDONLY)  
}
```

Change path

Mediation: TOCTTOU Vulnerabilities

Time-Of-Check-To-Time-Of-Use

```
void withdraw(int w){  
    b = getbalance();  
    if (b<w)  
        error("not enough $$");  
    b = b-w;  
    send(w)  
}
```



Buy

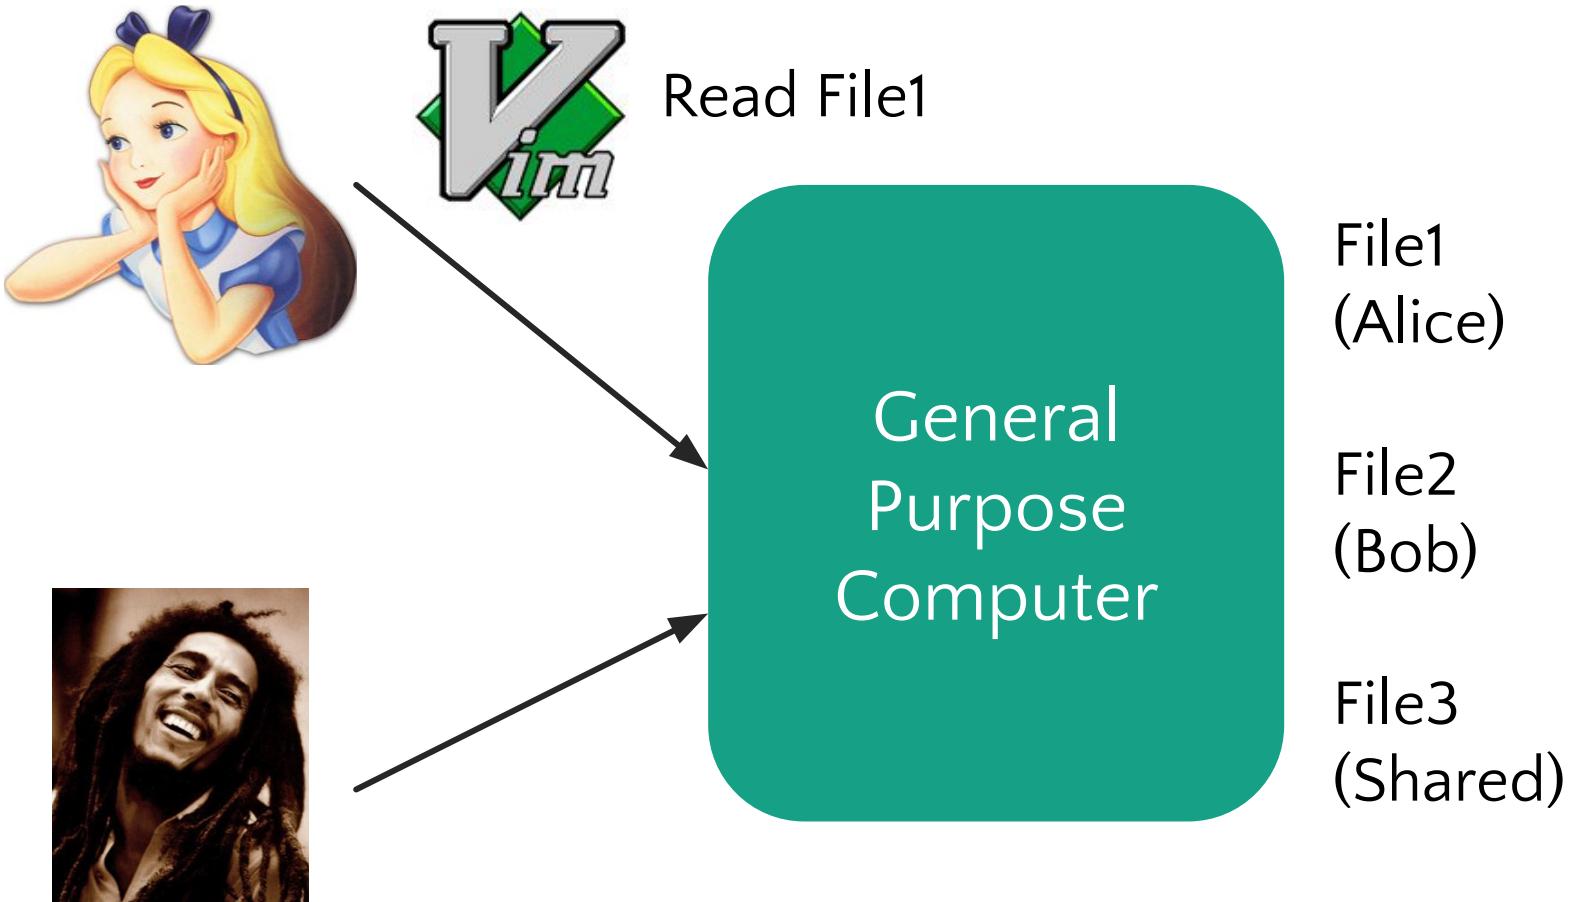
The Key Principles

- Economy of mechanism a.k.a. KISS
- Fail-safe defaults
- Don't rely on security by obscurity
- Complete mediation
- **Least privilege**
- Separation of duty
- Defense in depth
- Factor in users/acceptance/psychology
- Work factor/economics

} Later

Least Privilege

A user or entity should only have access to the specific data, resources and applications needed to complete a required task and nothing more.



The Key Principles

- Economy of mechanism a.k.a. KISS
- Fail-safe defaults
- Don't rely on security by obscurity
- Complete mediation
- Least privilege
- **Separation of duty**
- Defense in depth
- Factor in users/acceptance/psychology
- Work factor/economics

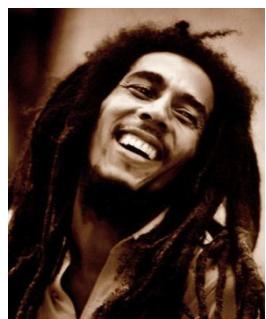
} Later

Separation of Duty (SoD)

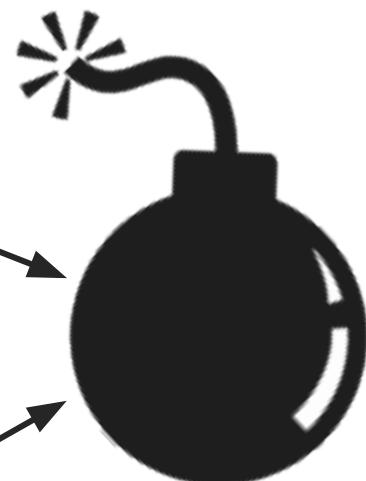
Having more than one person required to complete a task!



Push Pull Request
Launch Missile



Approve Pull Request
Launch Missile



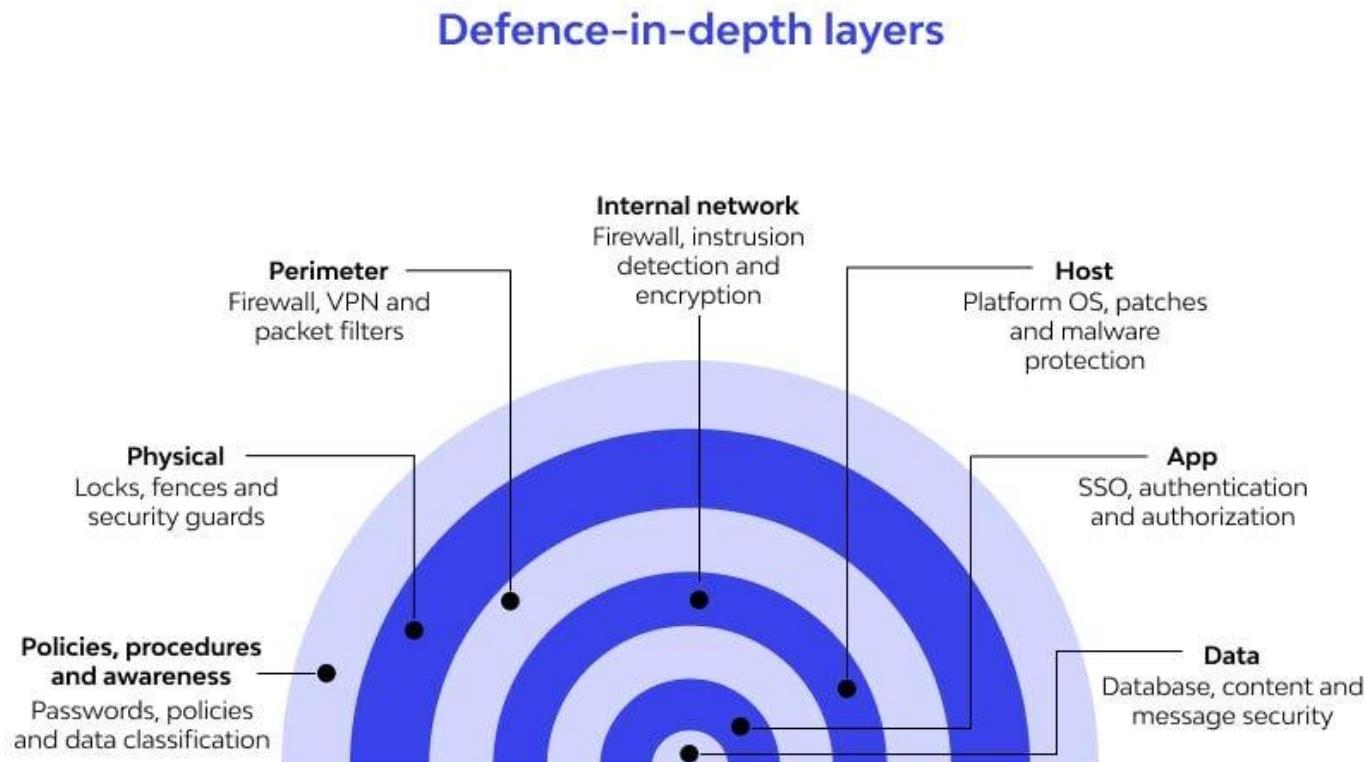
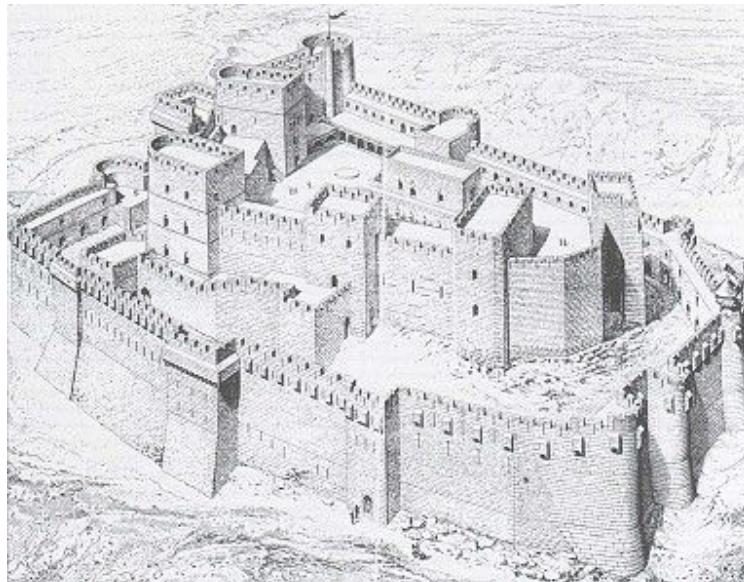
The Key Principles

- Economy of mechanism a.k.a. KISS
- Fail-safe defaults
- Don't rely on security by obscurity
- Complete mediation
- Least privilege
- Separation of duty
- **Defense in depth**
- Factor in users/acceptance/psychology
- Work factor/economics

} Later

Defense in Depth

- Few defensive measures are perfect
- Plan for failures
- Beware risk compensation!



Participation Question

Systems based on ACLs (like UNIX) typically deny access entirely if a subject is not listed on the relevant ACL. This is an example of which principle?

- A. Fail-safe defaults
- B. Complete mediation
- C. Separation of duty
- D. Defense in depth

I Followed All Security Principles - Am I Secure?



Ευχαριστώ και καλή μέρα εύχομαι!

Keep hacking!